

TW-Mailer Basic: Architecture and Implementation Report

Authors

- **Client Implementation:** Tarik Yilmaz
 - **Server Implementation:** Peyman Aparviz
-

Architecture

The TW-Mailer Basic system is designed as a **Client-Server application** using a simple text-based protocol over TCP.

- **Server:** Implements an **Iterative Server** model. It handles one client connection at a time using strict blocking I/O. It accepts a connection, processes all commands (`SEND`, `LIST`, `READ`, `DEL`) from that client in a loop until the client disconnects or sends `QUIT`, and then moves to the next waiting connection. This simplifies concurrency management by avoiding race conditions on the file system within a single process.
- **Client:** A command-line interface (CLI) that connects to the server and allows the user to interactively send commands. It manages user input parsing and protocol formatting.

Used Technologies

- **Language:** C++17.
- **Networking:** BSD/POSIX Sockets (`<sys/socket.h>`, `<netinet/in.h>`, `<arpa/inet.h>`) for reliable TCP communication.
- **File System:** strict usage of `std::filesystem` (C++17) for portable and robust file and directory manipulation (creating spool directories, listing files, removing messages).
- **Build System:** GNU Make.

Development Strategy

The development followed an agile, iterative approach:

- Shared Foundation:** A `common.hpp` header was created first to standardize constants (buffer sizes, command strings) and shared utility functions (`die`, `send_line`, `read_line`). This ensures protocol consistency between client and server.
- Client Implementation:** The client was implemented to validate the protocol format, user input handling, and strict socket communication.
- Server Implementation:** The server was built to handle the defined protocol, with a focus on robust file persistence.
- Refinement:** Enhancements like the "Auto-List" feature and timestamp display were added on top of the working core.

Implementation Details

Persistence Strategy: Timestamps and Sequence Numbers

To store messages safely and order them correctly without a complex database, we use a structured filename format:

`YYYYMMDD_HHMMSS_SEQ.txt` (e.g., `20251215_223005_001.txt`).

- **Why?** This format makes the filename **self-sorting**. An alphanumeric sort of the filenames typically corresponds to the chronological arrival order.
- **Sequence Number Check (Optimistic Write):** To determine the next available filename *without* scanning and counting all existing files (which would be slow $O(N)$), we use an optimistic check:
 1. Generate the base timestamp (e.g., `..._223005`).
 2. Start with sequence `001`.
 3. Check `if (fs::exists(filename_001))`.
 4. If it exists, increment (try `002`) and repeat. If not, this is our file.
 - *Benefit:* In the vast majority of cases, `001` is available. We only do multiple checks if high traffic occurs within the exact same second. This avoids the cost of iterating the entire directory for every write.

Reading and Ordering

While writing is optimized to avoid scanning, `LIST`, `READ`, and `DEL` operations require a consistent view of the mailbox.

- Since `fs::directory_iterator` does not guarantee any specific order, we iterate the directory, load all valid entries into a `std::vector`, and explicitly `std::sort` them.
- This ensures that "Message #1" is always the oldest message and "Message #N" is the newest, maintaining consistency between `LIST` indices and `READ` / `DEL` targets.

Date and Time Display

The `LIST` command enhances the user experience by extracting the date and time directly from the filename.

- The server parses the `YYYYMMDD_HHMMSS` portion of the filename.
- It formats this into a readable string (e.g., `[2025-12-15 22:30:05]`) and appends it to the subject line sent to the client.
- This allows the client to display timestamp information without needing to change the protocol structure (it simply prints the full subject line received).

Auto-List Feature

To improve usability and prevent errors, the Client tracks the "state" of the user's view.

- **Logic:** The client maintains a `last_list_user` variable.
- **Behavior:** When a user attempts to `READ` or `DEL` a message, the client checks if the current username matches `last_list_user`.
- **Automation:** If they differ (or if `LIST` was never called), the client automatically triggers a `LIST` command before sending the `READ` or `DEL` command. This ensures the user (and the logical index they provide) is synchronized with the latest server state for that specific mailbox.